

Amodei Answers His Critics, a Hacking AI Lands on Amazon, and Stripe's \$7B AI Bet

TL;DR

- ▶ Dario Amodei publicly rebutted a "world domination" accusation, reframing AI's PR problem as a trust crisis only real breakthroughs can fix.
- ▶ OpenAI's uncensored hacking model, GPT-5.6-Cyber, is now self-serve on Amazon's marketplace with 400+ known exploits already found.
- ▶ Stripe is buying AI marketplace OpenRouter for \$7B+, over 5x its valuation from three months ago.

AI TIPS

use AI better — how to do it

Fine-tune a specialist model instead of prompting a generalist

1. Install Unsloth Studio and pick an open-source model your hardware can run.
2. Build training examples (instruction + input + ideal output) — generate them with AI or convert a PDF into a dataset.
3. Train with QLoRA, which teaches a lightweight adapter instead of retraining the whole model.
4. Compare the fine-tuned model against the original to confirm your data actually improved it.

Source: The Neuron

Turn a chatbot into a cross-app Chief of Staff

1. Download an AI "Chief of Staff" agent, sign in, and start a focused workspace for one project.
2. Connect the apps it should touch — Gmail, Drive, Calendar, Notion, etc.
3. Ask it to map your workstreams and recommend automations to add, using the prompt below.
4. Have it produce a first real deliverable, then set up a daily debrief so you can see what it did each day.

Paste-ready: Review my connections, build a profile of my workstreams/projects, and suggest the three most useful agents, automations, and plugins I should add

Source: The Rundown AI

AI NEWS

the news — then what it means for you

Dario Amodei breaks his near-silence to rebut a "world domination" accusation. After investor Gavin Baker claimed Amodei privately believes Anthropic could become the only surviving private AI company, Amodei posted a two-part rebuttal that hit 10 million views in a day, calling public AI distrust "a crisis of trust" built on the industry's history of overpromising, not his own safety warnings.

What it means for you: A ready-made script for the trust conversation your clients are already having about AI vendors — Amodei's argument (real breakthroughs beat marketing) is one you can borrow when a client asks why they should believe any AI company's claims.

Source: The Neuron, The Rundown AI

OpenAI's uncensored hacking model is now on Amazon's cloud marketplace. GPT-5.6-Cyber, trained with fewer safety limits to write working exploit code, has already found over 400 privilege-escalation flaws in a single operating system — and no longer requires OpenAI's personal vetting to access.

What it means for you: If any client provisions AI tools through AWS, flag this now — self-serve exploit-writing AI shifts your threat model faster than most security reviews cycle.

Source: The Neuron

Stripe is buying AI model marketplace OpenRouter for over \$7 billion — more than 5x the \$1.3 billion valuation OpenRouter had just three months ago.

What it means for you: A signal that AI infrastructure is consolidating fast; if you're advising on vendor strategy, don't assume today's neutral marketplace stays neutral, or independent, for long.

AI anxiety is reshaping college majors. A Lumina Foundation/Gallup survey found 13–19% of degree students already changed their major over AI job fears, and computer science enrollment fell 8.1% last year after 15 years of growth, even as AI-specific majors jumped from 4 schools in 2020 to 79 today.

PM angle: Directly relevant — this is the incoming talent pipeline for every client hiring new grads; expect candidates to arrive with less CS depth and more AI-literacy framing on their resumes.

Conversation starter: *Are we seeing that shift yet in who's applying for entry-level roles, or is it still working its way through the pipeline?*

Source: Morning Brew

A Utah ice cream maker filed for bankruptcy after losing a \$23.8 million trademark lawsuit. A judge ruled Rebel Creamery intentionally copied rival Van Leeuwen's packaging; Rebel's bankruptcy filing lists \$13.8 million in assets against \$23.9 million in liabilities.

PM angle: Directly relevant — a clean, current example for any client conversation about brand or packaging IP risk, and how fast a copied look becomes an existential liability.

Conversation starter: *Worth a gut-check: how confident are we that our newest packaging or branding work wouldn't draw the same claim?*

Source: Morning Brew

Quant firm Jane Street lost \$15 billion in July in the fallout from the collapse of Situational Awareness, a \$45 billion AI-focused hedge fund run by 24-year-old Leopold Aschenbrenner.

PM angle: Directly relevant — a cautionary data point for clients treating an "AI" label as a proxy for "de-risked" when evaluating funds, vendors, or investments.

Conversation starter: *Has anyone actually asked what's inside the "AI" label before we treat it as a selling point?*

Source: 1440 Daily Digest

The US and Canada are racing to strike a trade deal before Wednesday's deadline. Without one, 50% tariffs hit \$20 billion of Canadian goods, from wine to cement, an outcome that could reshape the broader USMCA renegotiation.

PM angle: Directly relevant if any client sources from or sells into Canada — worth a proactive check on exposure before Wednesday.

Conversation starter: *Do we know which of our clients' supply chains actually touch Canadian imports right now?*

Source: Morning Brew

Sun-protective clothing is having a retail moment. UV-blocking apparel sales are up 12% year over year at US activewear brands, with searches for "UV shirt" up 26%, as celebrities like Anne Hathaway and Kendall Jenner popularize the look.

PM angle: No direct relevance to your work, general awareness only.

Conversation starter: *Anyone else noticing UV clothing everywhere this summer, or is that just my feed?*

Source: Morning Brew