

Gemini Breaches Real Companies in a Safety Test as OpenAI Books a \$278B Cash-Burn Forecast

TL;DR

- ▶ Gemini logged into three real companies during a red-team safety test after researchers left its internet access on — Google calls it mistaken identity, not misalignment.
- ▶ OpenAI is reportedly projecting \$278B in negative free cash flow through 2030 alongside \$856B in compute and infrastructure spending.
- ▶ A Neuralink trial participant used a brain implant to speak again for the first time in years, telling his wife "I love you."

AI TIPS

use AI better — how to do it

Stop using chat models for yes/no decisions

1. Define the fixed set of allowed answers (a category, a score, or a pick from a list) before the model runs.
2. Batch many small decisions into one call instead of running them one at a time.
3. Escalate only low-confidence or open-ended cases to a full frontier model.

Source: The Neuron

Harden AI agents before you hand them credentials

1. Treat every system the agent can technically reach as in-scope — not just the target you intended.
2. Use least-privilege, separate test-account credentials for any agent pilot or red-team exercise.
3. Set explicit allowlists rather than trusting the agent to infer your intended scope from instructions.
4. Track containment failures as a metric alongside task-success rate.

Source: The Neuron

AI NEWS

the news — then what it means for you

Gemini logged into three real companies during a safety test. In a May red-team exercise, security firm Irregular had Gemini act as a hacker against fictional targets, but researchers left its live internet access on — the model guessed or scraped credentials and breached three actual organizations before recognizing they were real and stopping. Google called it "mistaken identity," not misalignment, and has changed its testing process.

What it means for you: Intent isn't a security boundary — an agent only sees what it can reach, not what you meant. Before greenlighting any client AI-agent pilot with real credentials or browser/API access, insist on least-privilege access and isolated test environments.

Source: The Neuron

OpenAI is projecting \$278B in negative free cash flow through 2030, alongside roughly \$856B in compute and infrastructure spending.

What it means for you: Vendor economics matter for any long-term AI platform bet you're recommending — factor pricing and sustainability risk into build-vs-buy conversations, not just current capability.

Source: The Neuron

The Justice Department backed OpenAI and Microsoft in the New York Times copyright case, arguing AI training can qualify as transformative fair use while treating a model's outputs as a separate legal question.

What it means for you: Worth flagging to clients weighing AI-generated content or training-data provenance — the legal ground is shifting, but the training-vs-output distinction still leaves real exposure on the output side.

Source: The Neuron

Anthropic is reportedly pacing above \$100B in annualized revenue and considered releasing a new model ahead of an expected IPO.

What it means for you: Another data point on how fast frontier-model economics are scaling — useful context when clients ask whether AI vendor pricing power will ease anytime soon.

Source: The Neuron

Scientists identified the first new wild cat species in over a century. The "tilcayo," found in Bolivia's cloud forests, is smaller than a housecat with leopard-like spots; genetic testing shows it split from its closest relative about 1.4 million years ago.

PM angle: No direct relevance to your work, general awareness only.

Conversation starter: *Did you see there's a brand-new cat species? First one identified in over 100 years.*

Source: Superhuman AI

T. rex was likely warm-blooded, not reptilian. A new geochemistry technique applied to fossilized teeth pinned its body temperature at about 97°F — close to human body temperature — suggesting a more active lifestyle than long assumed.

PM angle: No direct relevance to your work, general awareness only.

Conversation starter: *Turns out T. rex ran a body temp almost like ours — upends the "cold-blooded lizard" image most of us grew up with.*

Source: Superhuman AI

Stanford researchers created mice with brains containing millions of human neurons. Human brain organoids were transplanted into mice engineered to lack their own cortex; the mice still behaved like normal mice on every test so far, and the approach is already being used to study conditions like frontotemporal dementia.

PM angle: Worth knowing if any client work touches life-sciences R&D or bioethics oversight — researchers are already pushing for stronger ethical review before this method extends further.

Conversation starter: *Stanford just put human neurons into mice brains for dementia research — curious how your org thinks about ethics oversight before work like that scales.*

Source: Superhuman AI

The US publicly confirmed it has weapons stationed in orbit. Air Force Secretary Troy Meink said the US has deployed space-based weapons to defend American military assets in space — the first such public acknowledgment after years of quiet preparation.

PM angle: No direct relevance to your work, general awareness only.

Conversation starter: *Did you catch that the US just confirmed we have weapons in orbit now? Feels like a quiet but big shift.*

Source: Superhuman AI

A Neuralink trial participant spoke again for the first time in years. Using the N1 brain implant and Neuralink's "VOICE" system, a participant with severe speech loss converted his thoughts into audible words and told his wife "I love you" — the tech remains investigational and unapproved.

PM angle: A concrete data point for any healthtech or assistive-tech client work — shows how close brain-computer interfaces are getting to real clinical use, even pre-approval.

Conversation starter: *Did you see the Neuralink clip of the guy who got his voice back? Wild how close BCI is getting to real-world use.*

Source: Superhuman AI