

OpenAI's Chief Scientist Warns of an Alignment Gap as Claude Cracks Fermat's Last Theorem

TL;DR

- ▶ OpenAI's chief scientist publicly called for an AI slowdown, warning no lab has solved alignment and monitoring well enough to keep scaling at full speed.
- ▶ Claude agents produced the first computer-verified proof of Fermat's Last Theorem in 11 days, a job mathematicians had budgeted years for.
- ▶ Nvidia's Jensen Huang declared "AGI has arrived," citing OpenAI's newest model training on 100,000 GPUs, a figure quietly revised down from an initial 300,000 claim.

AI TIPS

use AI better — how to do it

Run your own agent orchestration with just browser tabs

1. Planner pass: open a fresh chat, give it the full task, and ask it to break the work into 3-5 specific subtasks another AI could execute independently.
2. Worker passes: open a new chat per subtask and paste in only that subtask plus the context it needs, nothing from the other subtasks.
3. Verifier pass: open one more chat, paste in the original goal plus every worker's output, and ask it to catch contradictions, gaps, or quality issues before you ship it.

Paste-ready: Break this task into 3-5 specific, independent subtasks that a separate assistant could complete with no other context than what I give it: [YOUR TASK]

VERIFIER PROMPT:

Here's the original goal: [GOAL]

Here's what each subtask produced: [PASTE ALL OUTPUTS]

Check for contradictions, gaps, or quality issues before I ship this.

Source: The Neuron

Build a Lindy agent that never drops a follow-up

1. Create a Lindy account and connect one tool at a time, starting with Slack, Gmail, Calendar, Notion, or Drive instead of connecting everything.
2. Open Routines, then Discover, and pick a prebuilt workflow for meeting prep, follow-up, or weekly updates.
3. DM Lindy in Slack with a real task, then hand it a bigger assignment that can run in the background.
4. Add your phone number in Lindy so you can text your agent instead of opening the app.

Source: The Rundown AI

AI NEWS

the news — then what it means for you

OpenAI's chief scientist calls for an AI slowdown. Jakub Pachocki published an essay ("An Alien Mind") saying no lab has solved alignment and monitoring well enough to keep scaling responsibly; OpenAI's own data shows researchers now log 3.1 agent-workdays for every human workday (median researcher burning over \$600/day in inference), and the company is targeting a fully automated AI researcher by 2028.

What it means for you: This is a warning label from the company building the tech, not just outside critics -- worth citing when advising clients on how much human oversight to keep in the loop as they adopt agentic AI.

Source: The Neuron, The Rundown AI

Claude agents wrote the first computer-verified proof of Fermat's Last Theorem, done in 11 days. The proof ran to 13 million lines of code, a job mathematicians had budgeted years for.

What it means for you: A concrete data point for the "how far can agentic coding actually go" conversation with clients and leadership, even if it doesn't change your roadmap today.

Source: The Rundown AI

Nvidia's Jensen Huang declared "AGI has arrived," citing OpenAI's newest model (GPT-6 Astra) training on 100,000 Nvidia GPUs, with 400,000 more coming online next. Huang's post originally claimed 300,000 chips before being

deleted and quietly reposted with the smaller number, no explanation given.

What it means for you: Take the AGI framing with a grain of salt -- it's coming from OpenAI's chip supplier days after Nvidia posted \$89B in quarterly AI hardware sales -- but the compute scale-up itself is real and worth tracking.

Source: [The Neuron](#)

Microsoft rolled out Project Opal, a Copilot feature that hands off multi-step office tasks -- audit prep, IT ticket triage -- to an AI working on its own inside a virtual Windows PC. Currently limited to Frontier program testers.

What it means for you: Worth a pilot-program conversation with clients running Microsoft shops -- this is the kind of feature that will show up on Copilot roadmaps clients ask you about soon.

Source: [The Neuron](#)

Sources: [The Neuron](#), [The Rundown AI](#) (Sep 7). Newsletter links are click-tracked, so titles are cited rather than raw URLs.

impactmakers · AI Brief