

TypeSafe's Zero-Hallucination Model, Salesforce's In-House AI, and the Fed's First Hike in Three Years

TL;DR

- ▶ TypeSafe's new Jev model, from a ChatGPT co-creator, answers with confidence scores instead of prose — claimed zero hallucinations at \$42 per billion tokens, up to 200x faster than typical LLM workflows.
- ▶ Salesforce built its own in-house reasoning model, Koa, trained entirely on synthetic data — 3x fewer CRM errors than top models, with no customer data used.
- ▶ The Fed is expected to hike rates today for the first time in three years as Treasury yields hit 19-year highs — worth revisiting any client budgeting assumptions.

AI TIPS

use AI better — how to do it

Add confidence thresholds to your agent's decisions

1. Ask the agent for a fixed decision, not an essay.
2. Require a confidence score from 0 to 1 alongside the decision.
3. Set an escalation rule — route anything below a chosen threshold (e.g. 0.85) to a human or a stronger model.

Paste-ready: Classify this request as APPROVE, REVIEW, or REJECT. Return only the decision, a confidence score from 0 to 1, and one sentence explaining the most important uncertainty. If confidence is below 0.85, choose REVIEW.

Source: The Neuron

Give your coding agent access to cheaper models via OpenRouter

1. Sign in to OpenRouter, filter its catalog for your task, and sort low to high price.
2. Create a project-scoped API key with a spending limit and expiration date.
3. Give the key to your agent (Codex, Claude Code, etc.), run a small test task, then compare cost and quality before switching.

Source: The Rundown AI

AI NEWS

the news — then what it means for you

A ChatGPT co-creator's startup launched "Jev," an AI that answers with confidence scores instead of prose.

TypeSafe's Jev picks among pre-set options, responds in 70–500ms, and claims 40–200x faster, up to 238x cheaper results than typical LLM workflows, at \$42 per billion input tokens with free output.

What it means for you: A different tool category for narrow, high-volume decisions like routing, scoring, or screening — worth a look before defaulting to a full LLM for structured judgment calls.

Source: The Neuron, The Rundown AI

Salesforce built its own in-house reasoning model, Koa, trained entirely on synthetic data. Running on a customized Nvidia Nemotron 3 Super, Koa made 3x fewer errors than top models on Salesforce's internal CRM benchmark, with no customer data used.

What it means for you: Competitive intel if clients run on Salesforce — a build-vs-buy example of a vendor keeping AI and data in-house.

Source: The Rundown AI

OpenAI contractors have been reading real ChatGPT conversations to rate them, often without users' knowledge. A 404 Media investigation, "Project Lily," found hundreds of human reviewers sampling and grading live chats.

What it means for you: A concrete data point for the privacy question enterprise clients eventually ask about AI chat tools — worth having ready before they ask.

Source: The Rundown AI

The Fed is expected to raise interest rates today for the first time in three years. Chair Kevin Warsh cited a worse-than-expected inflation report and an oil price spike tied to the Iran war; markets are pricing a 92% chance of a 0.25% hike, and the 10-year Treasury yield hit 5.04% this week, its highest in 19 years.

PM angle: Revisit any client budgeting or financing conversations that assumed rates would hold or fall this quarter.

Conversation starter: *Are we still planning around steady rates, or does today's Fed move change the borrowing math for this project?*

Source: Morning Brew, 1440 Daily Digest

Cybersecurity stocks are surging as AI leaders warn their own tools could enable sophisticated hacks. CrowdStrike is up 17.6% and Palo Alto Networks up 13.5% this week, even as chip stocks fell; Gartner projects global cybersecurity spending will grow from \$51B this year to \$86B by 2027.

PM angle: A market signal worth citing when clients ask whether AI-driven security risk is real enough to budget for — investors are already voting with capital.

Conversation starter: *Have you noticed security spend is now tracking the same AI-risk headlines that used to just spook chip stocks?*

Source: Morning Brew

Canada pitched CAD\$1 trillion in investment projects to global asset managers, including a \$10 billion AI data center and a \$25 billion oil pipeline in Alberta, backed by a cut in effective tax rates on depreciating assets from 13% to 6.4%.

PM angle: A live example of a government using tax policy to court AI infrastructure investment — a useful comparison point if clients are weighing where to site data-heavy operations.

Conversation starter: *Are we tracking how aggressively Canada is courting AI infrastructure investment right now?*

Source: Morning Brew

Japan now has over 100,000 people age 100 or older for the first time, nearly 90% of them women, after 56 straight years of growth. The country needs an estimated 2.4 million elder-care workers in 2026, up from 2.1 million in 2024, and is nearly tripling foreign care-worker visas to fill the gap.

PM angle: A workforce-planning case study in aging demographics and immigration policy — relevant if any client work touches eldercare, healthcare staffing, or the Japan market.

Conversation starter: *If Japan's care workforce gap is this steep, what does that mean for any healthcare staffing clients eyeing that market?*

Source: Morning Brew